

***Finding Vulnerabilities in the Oracle VM
Virtualbox Hypervisor specific to CPU Attacks
and RAM Attacks***

Cyrus Majd

Honors Advanced Research

1/03/17

Project Abstract:

Cyber security plays a vital role in today's society and IT infrastructure. Nearly all of what we do eventually depends upon secure technologies. Virtualization claims to be the newest and most secure way to protect data and provide a secure cloud infrastructure, as they create "completely independent computer emulations". However, time and time again, malicious hackers have found vulnerabilities within hypervisors and have successfully stolen gigabytes of intellectual property from both companies and consumers alike. By exploiting and publishing key flaws in this technology before cyber thieves can, hundreds of potential threats will be eliminated.

My test procedures include stress testing both the computational and memory resources of two identical virtual machines running on the same hypervisor. My intention was to develop a methodical approach in a controlled environment in order to quantitatively determine the veracity of a cyber attack. This approach was split into two areas of study - CPU intensive attacks and RAM intensive attacks. Eventually, my results showed that both CPU and RAM intensive attacks are efficient when testing with one VM, but fail in affecting another through the hypervisor. These results are "in line" with previous studies on virtual machines, and prove that the hypervisor that I was using (Oracle VM Virtualbox) was indeed resistant to the CPU and RAM attacks I launched, although the single VM that I used was not.

Objectives:

Our three main objectives for this lab were to:

1. Observe whether or not virtual machines will break through the hypervisor (Oracle VM Virtualbox v.5.01) and share CPU power with one another when in stress when it achieves 100% utilization.
2. Observe whether or not virtual machines will break through the hypervisor (Oracle VM Virtualbox v.5.01) and share RAM storage with one another when the it is filled up.
3. Observe whether or not virtual machines will break through the hypervisor (Oracle VM Virtualbox v.5.01) and share hard drive storage with one another when their storage fills up.

Background:

An Introduction to Security

Cyber security's role in today's world is what governs most of our lives on the internet. Nearly all of what we do eventually depends upon secure technologies. It is the developer's job to make sure said technologies are safe and secure from cyber thieves and hackers, but even they make mistakes in their code; mistakes that allow hackers and cyber thieves access important files and financial data.

Cyber security is a field of study which can range from being a Certified Ethical Hacker to being a guidance counselor teaching students how to be safe on the internet. Being one of the hottest fields in the industry, cyber security experts face new challenges almost every day, whether it would be on developing new patches to faulty software, or discovering a new way a hacker could get into a system.

Virtualization is the act of partitioning the computer hardware and dedicating said partitions to a different OS (mostly) or app, creating independent virtual computers. Virtual

machines are disposable, meaning that they can be created and destroyed in software. Virtual machines can also act as proxies for large companies, and can be used as a cheap but reliable way to simulate real networks for students to study. Virtualization is NOT to be confused with hardware partitioning, since virtualization builds entirely independent compute machines. A hypervisor is the software that creates, destroys, and modifies virtual machines.

My project focuses on something a Certified Ethical Hacker would usually do; discover flaws in a specific software or system and publish them with a proposed solution. My project is about the security aspects of virtualization, specifically focusing on Oracle's virtualization technology for consumers called Virtualbox.

Identifying and publishing flaws in IT systems is very important because it helps developers to develop patches. With every new IT update there is a change to introduce security flaws or bugs. Hence certified ethical hackers and security analysts try to identify flaws in a system faster than hackers get a chance to exploit them. According to a NIST report from 2011 [6], if a "hypervisor provides a single point of security failure for all the guest Operating Systems; a single breach of the hypervisor places all the guest OSs at high risk." This just cannot be accepted in our technology centric life.

Advancements in Cyber Security

This topic has both a long and turbulent history. Ever since the internet was invented, hackers all over the world have been trying to misuse bugs and security flaws in systems, especially within virtualized cloud infrastructure, since many large businesses today rely on cloud services.

Ethical hacker organizations, such as Red Team Security, Offensive Security, CBT Nuggets, and the US Government all focus on virtualization security. However, most of their penetration tests are conducted on hypervisors developed by VMware and other prominent, expensive organizations. This leaves Oracle VM VirtualBox as a relatively untested hypervisor. For example, the last reported Nessus exploit report on VirtualBox was in January, 2016, which is VERY outdated for the IT world.

In addition, Oracle VM Virtualbox is much more different than VMware and other hypervisors because of its infrastructure on the host computer/server. There are essentially two different types of hypervisors, and the key difference between them is the way they are installed/configured. Below is a diagram depicting a common type 1 hypervisor.




Fig.1) A visual diagram of the infrastructure of a type 1 hypervisor. As shown, the type 1 hypervisor is the next direct layer past the host's firmware/BIOS

Type 1 hypervisors are installed directly on the firmware of the host machine, occupying the next layer beyond the firmware BIOS of the machine. Type 2 hypervisors are installed on top

of a compatible operating system, such as windows. Figure 2 demonstrates this.

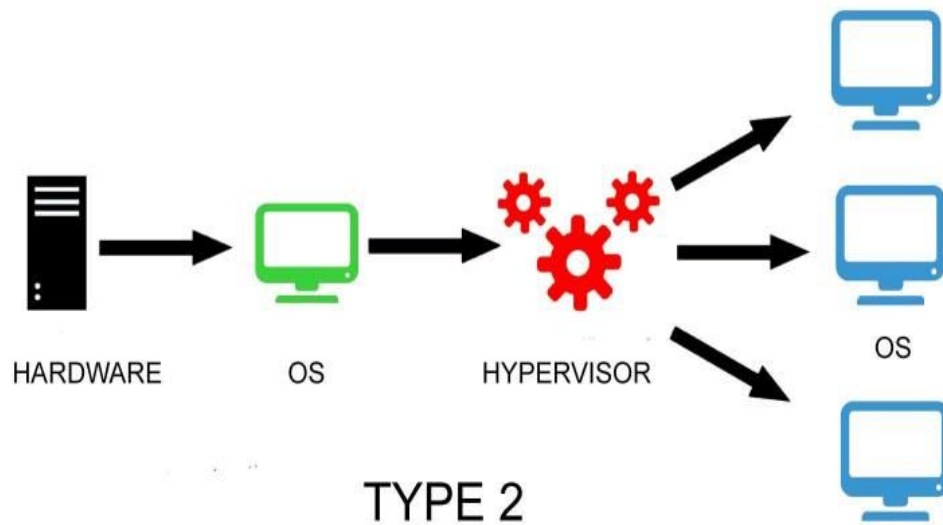


Fig 2. A diagram depicting a standard type 2 hypervisor.

Type 2 hypervisors are installed directly on top of the host OS. Oracle VM Virtualbox (the hypervisor that I will be working with) is an example of such a hypervisor.

The potential for program flaws increases dramatically with every new update. Below are a few examples of research conducted in cyber security and VMs.

Master Thesis from the University of Bergen, Norway, From R. Aarseth

A good example of such research was conducted by R. Aarseth from University of Bergen in Norway [9], who developed a comprehensive report on the security of Virtual Machines and their level of vulnerability to malwares such as Heartbleed. His results are used as reference in my research.

Aarseth began his experiments by first obtaining two different desktops, both on the same subnet. He reasons that the “timing attacks”, which according to Yinqian Zhang (a publisher of a

referenced article in Aarseth's works), makes it possible for hackers to steal encryption keys from VMs.

When he analyzed his data, Aarseth discovered that VMs may still be hacked if used on shared hardware, even if the VMs appear to be "isolated". According to his results, the methods used could potentially be "used to set up a more sophisticated attack against the users of a public cloud." He ends his report by theorizing that this may be a serious issue for almost all major public clouds.

Gartner's Report on The Security of Virtualized Servers

According to a Gartner Report from 2010 [12], 60 percent of virtualized servers will be less secure than the physical servers they replace by 2012. Gartner outlines in the report the six most common virtualization security risks and provides a recipe for how to avoid them. Gartner takes a different look at the cloud security and highlights that many of the cloud security flaws is caused because major virtualization deployment projects are being undertaken without involving the information security experts in the initial architecture and planning stages. Gardner states that virtualization is not inherently insecure. However, most virtualized workloads are being deployed insecurely, which is mainly a result of the immaturity of tools and processes and lack of good security experts.

International Journal of Computer Sciences paper on Cloud Computing Security Issues

Rabi Prasad et. al published a paper in the International Journal of Computer Science and IT Security [1], in which he outlines what cloud computing is and what are the various cloud

models and the main security risks and issues that are currently present within the cloud computing industry. This research paper analyzes the key research and challenges that are present in cloud computing and offers some of the best solutions to service providers and enterprises to avoid security failures.

The paper concludes that one of the biggest security worries with the cloud computing model is the sharing of resources. The authors discuss various models of cloud computing and security issues, and state the virtualization is major issue for cloud computing. The authors state that due to the complexity of the cloud, it will be difficult to achieve end-to-end security. They conclude that new security techniques need to be developed and older security techniques needed to be radically tweaked to be able to work with the clouds architecture.

The latest in cyber security research shows that there is a potential of threats in virtualization that may affect cyber security. Everybody agrees that the security of large cloud data centers is very important for our modern life, and with most datacenters running on virtualized platforms, it is very necessary to identify and eliminate flaws in different hypervisors.

Hypothesis:

1. A CPU attack on one of the VMs will slow down the processes of another identical VM.
2. A memory RAM attack of one VM will increase the memory utilization on another identical VM.
3. A CPU attack on one VM will fully consume all of the CPU power of that VM.
4. A RAM attack on one VM will fully consume all of the RAM power of that VM.

Experiment:

For the past year, I have been working on a security-based research project that involves discovering vulnerabilities(glitches, bugs, loopholes) in Ubuntu 16.04.1 and in Oracle VM

Virtualbox 15.01. I specifically was focusing on virtual machine “escapes”(when a VM user/file gains access to or is run by another VM without the permission of the hypervisor), and I figured that the best approach was to CPU stress test two identical virtual machines. I created two identical virtual machines, PC1 and PC2. The following are details on the two machines:

- a)Ubuntu 16.04.1 Release 64 Bit
- b)Base Memory/RAM-1024MB
- c)Boot order: Floppy Disk, Optical Drive, Hard Disk/SSD
- d)Video Memory-16MB
- e)Remote desktop server DISABLED
- f)Virtual hard disk encryption DISABLED
- g)1 processor dedicated, 1.6 GHz clock speed, 75% execution cap
- h)NAT network adapter ENABLED
- i)Virtual Hard Disk boot off UBUNTU ISO IMAGE
- j)Dynamically Allocated virtual storage system

The following is a list of materials used in this experiment:

- A. MacBook Air (host machine)
- B. Oracle VM Virtualbox 15.1 Release
- C. Ubuntu 16.04.1 Release 64 Bit
- D. TOP Ubuntu System Manager
- E. “count” shell script
- F. Memory attack shell script

I then coded a malicious script that makes the virtual machine run at 100% CPU. My plan was to stress test PC1 to see if it asks for some power that an idle PC2 has to spare. If PC2 was to donate such power-even a little- there is a vulnerability in the hypervisor, since I specifically completely isolated both virtual machines from the hypervisor.

The following is my attack code, “count”:

```
x="$1"

echo "Process $2 started at: " $(date +%s.%N)

while [ $x -gt 0 ];
do
    x=$(( x-1 ));
done

echo "process $2 ended at    " $(date +%s.%N)
```

As seen above, this file first creates a variable called x. It then prints out a message when the program starts and stops in nanoseconds since POSIX Time (January 1st, 1970, 00:00:00 AM).

The program then initiates a while loop that constantly repeats itself until the value of x (time in the input code below ie. 575000) is NOT greater than zero. It will then end the loop and print out when the program officially stopped.

On the first set of tests, program *count* was run at different niceness levels for in two instances with the following system inputs:

```
(bash count.sh 575000 A&);(nice -n 0 bash count.sh 575000 B&);sleep 1;ps
-l;sleep 20
```

The system input command is what runs the program (count) with specific parameters. The first part, *bash count*, addresses the program that is to be run. The number 575000 is the amount of time the program will run for (x in the “count” code). For this VM (PC1), 575000 translates to about 5-6 seconds of processing time, as intended. The time it takes to run this program differs from machine to machine, unless they both have the same **processing power** and general specs. The next part of the code, A&, simply defines a variable to this instance of the program. In this case, it is labeled as “A”. The next part serves the same purpose, and nice -n “0” is the command that sets the priority of the program, or “niceness”. In the above code, it is defaulted to zero. The rest of the input provides information on the CPU usage of the VM (ps -l) and also adds pause times to jot down notes (sleep 20, sleep 1). To run this code with a different niceness level, it must be typed in while logged in as an admin. This is achieved by typing in the command `sudo -i`, which lets you enter the “admin root”.

Data:

For niceness = 0, both instances of count (A and B) were found to have experienced roughly the same computing time (11.81 and 11.77 respectively) after comparing the averages of five tests. The data table shown below contains the results. (fig 1)

| | //Experimental Niceness: 0 | //Control Niceness: 0 |
|------------|-------------------------------|--------------------------|
| | B | A |
| start time | 1477927896.802263405 | 1477927896.807366127 |
| stop time | 1477927908.858708537 | 1477927909.126208482 |
| start time | 1477928511.441055882 | 1477928511.451852349 |
| stop time | 1477928523.220416999 | 1477928522.861984691 |
| start time | 1477928668.456747769 | 1477928668.470043854 |
| stop time | 1477928680.211829272 | 1477928680.438312757 |

| | | |
|--------------------|---|---|
| start time | 1477928854.772375868 | 1477928854.778360000 |
| stop time | 1477928866.506867748 | 1477928866.647110416 |
| start time | 1477929002.956806023 | 1477929002.973015639 |
| stop time | 1477929014.725855770 | 1477929014.261576803 |
| | <i>Difference is stop time - start time</i> | <i>Difference is stop time - start time</i> |
| Difference1 | 12.05644011 | 12.31884003 |
| Difference2 | 11.77936006 | 11.41013002 |
| Difference3 | 11.75507998 | 11.96826982 |
| Difference4 | 11.73448992 | 11.8687501 |
| Difference5 | 11.76905012 | 11.28855991 |
| | | |
| AVG of Differences | 11.81888404 | 11.77090998 |

(fig 1)

The CPU distributes the tasks and divides the computing time in half. This is why the intended 6 second long programs now each take about 11-12 seconds to complete. The CPU does this because the two programs share the same niceness index and have the same code, so they are equally as demanding for resources. Next, the two processes were tested with different niceness levels. Our control group process (A) remained with a niceness of 0 (default) but our experimental group process (B) had its niceness level decreased to -5 in order to (theoretically) consume more CPU (and thus have a shorter runtime) than our control group. The results backed up this hypothesis.

| | |
|--------------------------------|--------------------------|
| | |
| //Experimental Niceness: -5 | //Control Niceness: 0 |
| B | A |
| 1478013548.253796412 | 1478013548.274454900 |
| 1478013556.335801673 | 1478013560.078921220 |
| 1478013734.315090175 | 1478013734.321089798 |

| | |
|--------------------------------|--------------------------------|
| 1478013742.155983718 | 1478013746.101238516 |
| 1478013860.545381114 | 1478013860.573184259 |
| 1478013868.536137910 | 1478013872.110450959 |
| 1478014034.020826585 | 1478014034.027257412 |
| 1478014042.054679765 | 1478014045.947034306 |
| 1478018113.729885014 | 1478018113.757396859 |
| 1478018121.276924735 | 1478018124.916611756 |
| Differences: | Differences: |
| 8.082010031 | 11.80446982 |
| 7.840890169 | 11.78014994 |
| 7.990749836 | 11.53727007 |
| 8.033850193 | 11.91978002 |
| 7.547039986 | 11.15921998 |
| Average of Differences: | Average of Differences: |
| 7.898908043 | 11.64017797 |

(fig 2)

As shown, the time it took for the experimental group process to compute decreased by about 4 seconds, whereas the control group process took the same amount of time. The programs are still both running at the same time, but now one instance finishes faster than another. The CPU usage remains at around 98-100% the entire time.

The following are the results of testing with an experimental group niceness of -10.(fig 3)

| //Experimental Niceness: -10 | //Control Niceness: 0 |
|---------------------------------|--------------------------|
| B | A |
| 1478108355.761422928 | 1478108355.842813791 |
| 1478108362.608876104 | 1478108367.880254587 |
| 1478108629.938161597 | 1478108630.036205762 |
| 1478108637.711882039 | 1478108642.684697574 |
| 1478108943.285563098 | 1478108943.323189231 |
| 1478108950.342262971 | 1478108955.276416667 |

| | |
|--------------------------------|--------------------------------|
| 1478109093.890490374 | 1478109093.938032693 |
| 1478109100.372786445 | 1478109105.655309886 |
| 1478109228.167464276 | 1478109228.200227203 |
| 1478109234.698023888 | 1478109239.706137542 |
| Differences: | Differences: |
| 6.847450018 | 12.03744006 |
| 7.773720026 | 12.64848995 |
| 7.056699991 | 11.95323014 |
| 6.48229003 | 11.7172699 |
| 6.530560017 | 11.50590992 |
| Average of Differences: | Average of Differences: |
| 6.938144016 | 11.97246799 |

(fig 3)

Something very interesting happens with this specific configuration. At the niceness of 0, both processes took about 12 seconds to complete, as previously mentioned. The code tested was identical for both of the processes, so both process started and stopped at the same time, but with only 50% of the CPU's computing resources for each(for a total of 100% CPU usage). At the niceness of -5, both processes were running simultaneously, but experimental group process B consumes more CPU power than when tested with a niceness of 0. The total amount of CPU usage was still 100%, as expected. However, when process B is assigned to a niceness of -10, it finished about five seconds before process A. Because it is known that process A can run at a maximum of five seconds while given 100% CPU, and because process A started at the same time as process B, it can be deducted that process B consumed ~100% CPU for 5-6 seconds to complete it's own process before handing over the full CPU power to process A. **In other words, process B seems to have prevented process A from computing anything before it**

completed itself. This is a discovery by itself, proving that it is very possible to induce a denial of service attack on this specific version of Ubuntu when logged in as a sudo root user.

As expected, the following tests conducted with experimental group niceness values of -15 and -20 respectively display very similar results.

| //Experimental Niceness: -15 | //Control Niceness: 0 |
|---------------------------------|--------------------------------|
| B | A |
| 1478893218.328216650 | 1478893218.321407109 |
| 1478893223.942201169 | 1478893229.228214130 |
| 1478893672.413309477 | 1478893672.405256531 |
| 1478893678.384173091 | 1478893683.992002030 |
| 1478893840.332820033 | 1478893840.331578030 |
| 1478893846.461353992 | 1478893851.817575111 |
| 1478892961.886865837 | 1478893961.878954483 |
| 1478892967.717295756 | 1478893972.931884809 |
| 1478894174.517993602 | 1478894174.511188082 |
| 1478894180.375599563 | 1478894185.954847456 |
| Differences: | Differences: |
| 5.61398983 | 10.90681005 |
| 5.970870018 | 11.58675003 |
| 6.128530025 | 11.48600006 |
| 5.830430031 | 11.05292988 |
| 5.857599974 | 11.44366002 |
| Average of Differences: | Average of Differences: |
| 5.880283976 | 11.29523001 |

(fig 4)

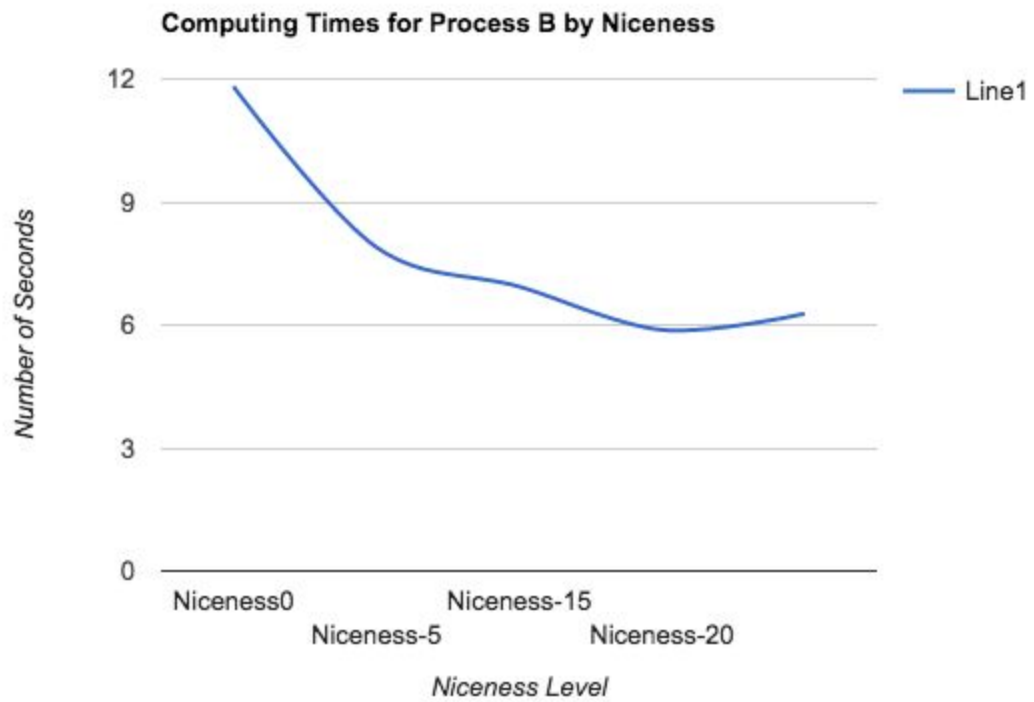
| //Experimental Niceness: -20 | //Control Niceness: 0 |
|---------------------------------|--------------------------|
| B | A |
| 1479085662.716082866 | 1479085662.698009499 |

| | |
|--------------------------------|--------------------------------|
| 1479085668.198122511 | 1479085673.245204646 |
| 1479085748.767270061 | 1479085748.760568797 |
| 1479085755.279890730 | 1479085760.919323796 |
| 1479085807.016137317 | 1479085807.003721149 |
| 1479085813.357224751 | 1479085819.408789059 |
| 1479085856.558462184 | 1479085856.557135585 |
| 1479085863.019615272 | 1479085869.223539058 |
| 1479085896.198018431 | 1479085896.190087347 |
| 1479085902.768550087 | 1479085908.484534021 |
| Differences: | Differences: |
| 5.482040167 | 10.54719996 |
| 6.512619972 | 12.15876007 |
| 6.341089964 | 12.40506005 |
| 6.461149931 | 12.66639996 |
| 6.570539951 | 12.29445004 |
| Average of Differences: | Average of Differences: |
| 6.273487997 | 12.01437402 |

(fig 5)

We can see from the average of the differences in the above two tables (figs 4 & 5) that results are nearly identical to the experiment with an experimental niceness of -10. It seems as if the niceness of 10 is the cutoff line that separates the sharing of CPU resources and the hoarding of it.

The following graph demonstrates how the computing time slowly decreases and levels out. (fig 6)




(fig 6)

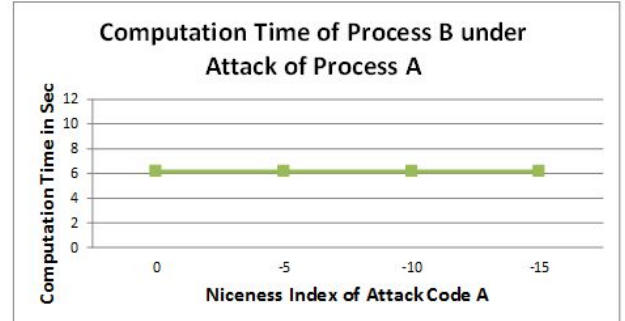
This next graph shows in total the results derived from this test, along with the results from the next test. In the graph, the attack code is called process A and the user test code is called process B. (fig 7)

CPU-Intensive Attack

Run on the same VM



Run on different VMs



The "Control Group" for CPU-Attack Experiments:

- Attack Code A and Process B are two separate software, each one lasting for 5 seconds when run alone.
- A and B run on the same VM, sharing the CPU.
- A slows down B as the Niceness Index decreases
- At Niceness Index of -20, A freezes B for 5 seconds.
- The experiment verifies the "Operating System Leak" caused by attack code A. The attack spills over to B.

The "Experimental Group" for CPU-Attack Experiments:

- Attack Code A runs continuously on VM1 driving its CPU utilization to ~100%
- Process B runs on VM2, while its Computation Time is measured
- A does not slow down B (independent of the Niceness Index)
- A and B run on different VMs, but share the hypervisor. It means, A and B are running on the same cloud.
- No hypervisor leak is detected. The attack does not spill over to B.

(fig 7 above)

The second series of tests were conducted while monitoring two VMs at the same time (PC1 and PC2). The plan was to run count on PC2 for an infinite amount of time(forever) while then running count on PC1 at different niceness levels. Crossovers, if there were any, would be indicated by a computation time on PC1 other than 5-6 seconds. Fig. 8 shows the results for the tests with the niceness at 0.

| | PC1 | PC1 | PC2 |
|------------|----------------------|---------------|---------------|
| | //Experimental | //Utilization | //Utilization |
| | Niceness: 0 | CPU | CPU |
| | | | |
| start time | 1479089728.535449219 | 94% | 93% |
| stop time | 1479089734.001340316 | 94% | 93% |
| start time | 1479089843.050117253 | 93% | 93% |

| | | | |
|--------------------|--------------------------------|-----|-----|
| stop time | 1479089848.445491476 | 93% | 93% |
| start time | 1479089881.679132445 | 92% | 93% |
| stop time | 1479089887.000384434 | 92% | 93% |
| start time | 1479089919.288704391 | 94% | 93% |
| stop time | 1479089924.670074507 | 94% | 93% |
| start time | 1479089956.477907653 | 94% | 93% |
| stop time | 1479089961.810057752 | 94% | 93% |
| | Differences: | | |
| Difference1 | 5.465899944 | | |
| Difference2 | 5.395379782 | | |
| Difference3 | 5.321249962 | | |
| Difference4 | 5.381369829 | | |
| Difference5 | 5.332149982 | | |
| | Average of Differences: | | |
| AVG of Differences | 5.3792099 | | |

(fig 8)

Unfortunately, no significant difference in computation time was found. I wasn't very lucky with the rest of the niceness tests either, as seen in figures 9, 10, 11, and 12.

| PC1 | PC1 | PC2 |
|--------------------------------|----------------------|----------------------|
| //Experimental Niceness: -5 | //Utilization CPU | //Utilization CPU |
| | | |
| 1479090521.794296757 | 96% | 95% |
| 1479090527.175383264 | 96% | 95% |
| 1479090569.245218864 | 94% | 95% |
| 1479090574.533528898 | 94% | 95% |
| 1479090606.078422123 | 95% | 95% |
| 1479090611.365465988 | 95% | 95% |
| 1479090637.704672527 | 93% | 95% |
| 1479090642.929023125 | 93% | 95% |
| 1479090677.938510570 | 93% | 95% |

| | | |
|--------------------------------|-----|-----|
| 1479090683.302848625 | 93% | 95% |
| Differences: | | |
| 5.381089926 | | |
| 5.288310051 | | |
| 5.287039995 | | |
| 5.224349976 | | |
| 5.364330053 | | |
| Average of Differences: | | |
| 5.309024 | | |

(fig 9 above)

| PC1 | PC1 | PC2 |
|--------------------------------|---------------|---------------|
| //Experimental | //Utilization | //Utilization |
| Niceness: -10 | CPU | CPU |
| | | |
| 1479144699.290128550 | 98% | 99% |
| 1479144704.597686437 | 98% | 99% |
| 1479144754.766869313 | 94% | 99% |
| 1479144759.947665585 | 94% | 99% |
| 1479144790.096480499 | 96% | 99% |
| 1479144795.543738782 | 96% | 99% |
| 1479144823.554442141 | 92% | 99% |
| 1479144828.758479108 | 92% | 99% |
| 1479144859.242001610 | 96% | 99% |
| 1479144864.477946130 | 96% | 99% |
| Differences: | | |
| 5.307560205 | | |
| 5.180799961 | | |
| 5.447250128 | | |
| 5.204030037 | | |
| 5.23593998 | | |
| Average of Differences: | | |
| 5.275116062 | | |

(fig 10 above)

| PC1 | PC1 | PC2 |
|---------------------------------|----------------------|----------------------|
| //Experimental Niceness: -15 | //Utilization CPU | //Utilization CPU |
| | | |
| 1479146058.287559862 | 94% | 99% |
| 1479146063.646615119 | 94% | 99% |
| 1479146120.573801707 | 94% | 99% |
| 1479146126.770548664 | 94% | 99% |
| 1479146152.427704998 | 94% | 99% |
| 1479146158.131679942 | 94% | 99% |
| 1479146179.366264436 | 94% | 99% |
| 1479146185.247508683 | 94% | 99% |
| 1479146233.221250231 | 95% | 99% |
| 1479146238.549494526 | 95% | 99% |
| Differences: | | |
| 5.359060049 | | |
| 6.196739912 | | |
| 5.703969955 | | |
| 5.881239891 | | |
| 5.328239918 | | |
| Average of Differences: | | |
| 5.693849945 | | |

(fig 11 above)

| PC1 | PC1 | PC2 |
|---------------------------------|----------------------|----------------------|
| //Experimental Niceness: -20 | //Utilization CPU | //Utilization CPU |
| | | |
| 1479235617.728908050 | 94% | 99% |
| 1479235623.357723942 | 94% | 99% |
| 1479235690.038594346 | 97% | 99% |

| | | |
|--------------------------------|-----|-----|
| 1479235695.744354110 | 97% | 99% |
| 1479235725.523000376 | 92% | 99% |
| 1479235730.952676244 | 92% | 99% |
| 1479235763.329573481 | 94% | 99% |
| 1479235768.673714364 | 94% | 99% |
| 1479235795.152635043 | 96% | 99% |
| 1479235800.532483973 | 96% | 99% |
| Differences: | | |
| 5.628819942 | | |
| 5.705760002 | | |
| 5.429670095 | | |
| 5.344140053 | | |
| 5.379849911 | | |
| Average of Differences: | | |
| 5.497648001 | | |

(fig 12 above)

At this point I have tested to see if CPU intensive attacks can trigger a virtual machine escape from Oracle VM Virtualbox and if it is possible to perform a denial of service attack against a user with a CPU intensive attack with different niceness levels. Although my hypothesis that the CPU attack would create a virtual machine escape was proven incorrect, my denial of service attack undoubtedly did work as intended. These discoveries can be published as both new data (new method of denial of service attack) and supportive data for other projects (virtual machine escapes are impossible in Oracle VM Virtualbox 5.1 with CPU intensive programs). The graph on the right hand side of figure 7 depicts the results of this test. Process B is portrayed as the passive user application while process A is shown as the attack malware.

The next step in my experiments was to measure any RAM overflow from either virtual machine. To do this, I first had to develop a special software that would be designed to overflow the target virtual machine's RAM buffer. I named the program "m-attack.sh", short for "memory attack bash file". The following is the source code of m-attack.sh:

```
#!/bin/bash

echo "Process m-attack started at: " $(date +%s.%N)

x="$1";
y="01234567";
while [ $x -gt 0 ];
do
    x=$(( x-1 ));
    z=$y$y;
    y=$z;
done
while [ TRUE ];
do
    z=$y;
    y=$z;
done
```

As shown above, the program's essence lies within its two vital "for loops" after receiving a user-input for variable "x" and another input for y. Variable x goes through a for loop that repeats itself as long as x is greater than zero, and with every iteration, variable y is doubled and

placed in a temporary variable z, which is then returned as the new y value after the second while loop is completed.

The charts below show the data collected while running both m-attack and count on the same machine. Note that the count process was run with the default priority level, since the test was intended to mainly analyze m-attack.

| | | | | | | |
|------------|------------------------------|-------------------------------|---------------|----------------|---------------------|----------|
| | PC1 | PC1 | PC1 | PC1 | PC1 | PC1 |
| | m-attack | count | max memory | used memory | mem-utiliza tion | CPU util |
| | | | | | | |
| | #Of m-attack instances | time measurements of count | | | | |
| start time | 1 | 1483217249.97173020 8 | 1016252 | 547088 | 53.83% | 99% |
| stop time | 1 | 1483217261.19638873 4 | | | | |
| start time | 1 | 1483221813.88294572 2 | 1016252 | 520156 | 51.18% | 98.8% |
| stop time | 1 | 1483221825.68351768 8 | | | | |

| | | | | | | | |
|-------------|---|---------------------|---|---------|-----------------|--------|-----|
| start time | 1 | 1483222023.49180033 | 2 | 1016252 | 668532 | 65.78% | 99% |
| stop time | 1 | 1483222034.58174736 | 6 | | | | |
| start time | 1 | 1483222882.74908763 | 3 | 1016252 | 665512 | 65.49% | 99% |
| stop time | 1 | 1483222893.80292167 | 7 | | | | |
| start time | 1 | 1483222977.76981081 | 0 | 1016252 | 618860 | 60.90% | 98% |
| stop time | 1 | 1483222989.94230768 | 5 | | | | |
| | | | | | | | |
| Difference1 | | 11.22464991 | | | AVG mem use: | 59.44% | |
| Difference2 | | 11.80057001 | | | | | |
| Difference3 | | 11.08993983 | | | | | |
| Difference4 | | 11.05384016 | | | | | |
| Difference5 | | 12.17249012 | | | | | |

(figure 13 above)

As seen above, running one instance of m-attack along with count takes up about 59.44% of total memory space, while constantly using about 99% of CPU resources. The process time for Count was around 11-12 seconds, which was expected.

| | | | | | | |
|------------|------------------------------|-------------------------------|---------------|----------------|---------------------|----------|
| | PC1 | PC1 | PC1 | PC1 | PC1 | PC1 |
| | m-attack | count | max memory | used memory | mem-utiliza tion | CPU util |
| | | | | | | |
| | #of m-attack instances | time measurements of count | | | | |
| start time | 2 | 1483217455.41293720 4 | 1016252 | 833324 | 82.00% | 99% |
| stop time | 2 | 1483217469.17616308 8 | | | | |
| start time | 2 | 1483223148.33764692 4 | 1016252 | 791140 | 77.85% | 99% |
| stop time | 2 | 1483223158.75882373 4 | | | | |
| start time | 2 | 1483223296.80767808 6 | 1016252 | 801140 | 78.83% | 98.4% |
| stop time | 2 | 1483223307.68161906 9 | | | | |
| start time | 2 | 1483223406.41114214 2 | 1016252 | 845656 | 83.21% | 99% |

| | | | | | | |
|-------------|---|--------------------------|---------|-----------------|--------|-----|
| stop time | 2 | 1483223417.50841531 9 | | | | |
| start time | 2 | 1483223504.53605193 7 | 1016252 | 844488 | 83.10% | 98% |
| stop time | 2 | 1483223515.31671480 6 | | | | |
| | | | | | | |
| Difference1 | | 13.76323009 | | AVG mem use: | 81.00% | |
| Difference2 | | 10.42118001 | | | | |
| Difference3 | | 10.87393999 | | | | |
| Difference4 | | 11.09727001 | | | | |
| Difference5 | | 10.78065991 | | | | |

(fig 14 above)

In this test, two parallel instances of m-attack were launched in the background, but count was still called only once. This time, the average memory utilization increased to about 81%, which was significantly higher than the previous 59.44% utilization.

| | | | | | | |
|--|----------|-------|---------------|----------------|---------------------|----------|
| | PC1 | PC1 | PC1 | PC1 | PC1 | PC1 |
| | m-attack | count | max memory | used memory | mem-utiliza tion | CPU util |
| | | | | | | |

| | #Of m-attack instances | time measurements of count | | | | |
|------------|------------------------------|-------------------------------|---------|--------|--------|-------|
| start time | 3 | 1483217841.69157367 4 | 1016252 | 871940 | 85.80% | 99% |
| stop time | 3 | 1483217854.56928873 7 | | | | |
| start time | 3 | 1483218069.19785843 1 | 1016252 | 874228 | 86.02% | 98.5% |
| stop time | 3 | 1483218081.64473672 5 | | | | |
| start time | 3 | 1483218200.87872123 6 | 1016252 | 861060 | 84.73% | 98.7% |
| stop time | 3 | 1483218213.22373759 7 | | | | |
| start time | 3 | 1483218348.42416039 1 | 1016252 | 861436 | 84.77% | 99% |
| stop time | 3 | 1483218359.09593955 3 | | | | |
| start time | 3 | 1483218468.86225632 1 | 1016252 | 861452 | 84.77% | 98% |
| stop time | 3 | 1483218480.97560914 1 | | | | |
| | | | | | | |

| | | | | | | |
|-------------|--|-------------|--|-----------------|--------|--|
| Difference1 | | 12.87770987 | | AVG mem use: | 85.22% | |
| Difference2 | | 12.4468801 | | | | |
| Difference3 | | 12.34501004 | | | | |
| Difference4 | | 10.6717701 | | | | |
| Difference5 | | 12.11334991 | | | | |

| | | | | | | |
|------------|------------------------------|-------------------------------|---------------|----------------|---------------------|----------|
| | PC1 | PC1 | PC1 | PC1 | PC1 | PC1 |
| | m-attack | count | max memory | used memory | mem-utiliza tion | CPU util |
| | | | | | | |
| | #of m-attack instances | time measurements of count | | | | |
| start time | 4 | 1483218675.49351394 3 | 1016252 | 845240 | 83.17% | 99% |
| stop time | 4 | 1483218687.69813471 9 | | | | |
| start time | 4 | 1483218866.35634192 8 | 1016252 | 844188 | 83.07% | 99% |
| stop time | 4 | 1483218877.95042928 5 | | | | |

| | | | | | | |
|-------------|---|--------------------------|---------|-----------------|--------|-----|
| start time | 4 | 1483218998.59438992 2 | 1016252 | 843964 | 83.05% | 99% |
| stop time | 4 | 1483219009.58468997 3 | | | | |
| start time | 4 | 1483221365.72977034 0 | 1016252 | 843932 | 83.04% | 99% |
| stop time | 4 | 1483221377.99755752 3 | | | | |
| start time | 4 | 1483221514.07287474 7 | 1016252 | 847872 | 83.43% | 98% |
| stop time | 4 | 1483221525.96086813 3 | | | | |
| | | | | | | |
| Difference1 | | 12.20461988 | | AVG mem use: | 83.15% | |
| Difference2 | | 11.59407997 | | | | |
| Difference3 | | 10.99030018 | | | | |
| Difference4 | | 12.26778007 | | | | |
| Difference5 | | 11.88799 | | | | |

The following charts show M-Attack's results after running on PC1 while Count was running on PC2. The number of instances of M-Attack is expected to affect the process time of Count.

| | PC1 | PC2 | PC1 | PC1 | PC1 |
|------------|----------|--------------------------|------------|-------------|-----------------|
| | m-attack | count | max memory | used memory | mem-utilization |
| | | | | | |
| | | | | | |
| start time | 1 | 1483124777.57836180 5 | 1016252 | 812708 | 79.97% |
| stop time | 1 | 1483124783.26385501 9 | | | |
| start time | 1 | 1483125234.14125905 4 | 1016252 | 765988 | 75.37% |
| stop time | 1 | 1483125240.02819630 5 | | | |
| start time | 1 | 1483125312.50640396 1 | 1016252 | 784524 | 77.20% |
| stop time | 1 | 1483125318.30457777 0 | | | |
| start time | 1 | 1483125344.37517569 8 | 1016252 | 777108 | 76.47% |

| | | | | | |
|-------------|---|--------------------------|---------|--------|--------------|
| stop time | 1 | 1483125350.49084902 3 | | | |
| start time | 1 | 1483125514.52250339 6 | 1016252 | 769984 | 75.77% |
| stop time | 1 | 1483125520.37435969 9 | | | |
| | | | | | |
| Difference1 | 0 | 5.685489893 | | | 0.7695555827 |
| Difference2 | 0 | 5.886940002 | | | |
| Difference3 | 0 | 5.798169851 | | | |
| Difference4 | 0 | 6.115669966 | | | |
| Difference5 | 0 | 5.851850033 | | | |

(Above: M-Attack ran on PC1 (one instance) while Count runs on PC2)

As shown above, the first m-attack test resulted in an unusually high memory utilization for PC1. The average memory utilization is 77%, and the average computing time for Count on PC2 is about 5.867623949 seconds.

| | | | | |
|----------|--------------------------|------------|-------------|-----------------|
| PC1 | PC2 | PC1 | PC1 | PC1 |
| m-attack | count | max memory | used memory | mem-utilization |
| | | | | |
| | | | | |
| 2 | 1483125689.25204921 4 | 1016252 | 813968 | 80.10% |
| 2 | 1483125694.94848451 4 | | | |

| | | | | |
|--------------|----------------------|---------|--------|--------------|
| 2 | 1483125731.093337276 | 1016252 | 817068 | 80.40% |
| 2 | 1483125737.023744375 | | | |
| 2 | 1483125811.091840130 | 1016252 | 817068 | 80.40% |
| 2 | 1483125816.740025071 | | | |
| 2 | 1483125925.386134430 | 1016252 | 816340 | 80.33% |
| 2 | 1483125931.133659217 | | | |
| 2 | 1483126103.169074872 | 1016252 | 812536 | 79.95% |
| 2 | 1483126109.215024679 | | | |
| | | | | |
| Difference 1 | 5.696439981 | | | 0.8023561085 |
| Difference 2 | 5.930410147 | | | |
| Difference 3 | 5.648180008 | | | |
| Difference 4 | 5.74751997 | | | |
| Difference 5 | 6.045949936 | | | |

(Above: M-Attack ran on PC1 (two instances) while Count runs on PC2)

This second test seems to yield more consistent data points on memory utilization in particular. All five data points closely revolve around 80%, which is nearly a 3% memory utilization increase from the previous test. This deviation is relatively small, so there is effectively no profound change in PC2's memory utilization. Count's process time is about 5.813700008 seconds.

| PC1 | PC2 | PC1 | PC1 | PC1 |
|----------|-------|------------|-------------|-----------------|
| m-attack | count | max memory | used memory | mem-utilization |
| | | | | |
| | | | | |


| | | | | |
|--------------|--------------------------|---------|--------|--------------|
| 3 | 1483128566.03368673 5 | 1016252 | 859904 | 84.62% |
| 3 | 1483128571.75336391 2 | | | |
| 3 | 1483129196.05336302 4 | 1016252 | 761500 | 74.93% |
| 3 | 1483129201.99217558 2 | | | |
| 3 | 1483129419.22986978 4 | 1016252 | 762936 | 75.07% |
| 3 | 1483129425.01526510 1 | | | |
| 3 | 1483129494.16129174 1 | 1016252 | 775148 | 76.28% |
| 3 | 1483129499.94125281 0 | | | |
| 3 | 1483129551.42164132 5 | 1016252 | 764972 | 75.27% |
| 3 | 1483129557.02379012 0 | | | |
| | | | | |
| Difference 1 | 5.719680071 | | | 0.7723399314 |
| Difference 2 | 5.93881011 | | | |
| Difference 3 | 5.785399914 | 1016272 | 856100 | 84.24% |
| Difference 4 | 5.779960155 | 1016272 | 778912 | 76.64% |
| Difference 5 | 5.602149963 | 1016272 | 788660 | 77.60% |

(Above: M-Attack ran on PC1 (three instances) while Count runs on PC2)

This test above runs three instances of M-Attack on PC1 while an instance of Count runs on PC2. In this test, the average memory utilization of PC1 hits about 77.23%, and the average process time for Count is about 5.7652000426 seconds. The memory utilization never really experienced a significant change in these series of tests, which implies that there is no leak between VMs while running my memory attack code.


The results from all of the RAM Attacks are seen here, in figure 15.

Memory-intensive Attack



The "Control Group" for Memory-Attack Experiments:

- Attack Code A runs continuously on VM1, gradually filling up memory
- Process B fills up the memory in a range of ~60% - 85%
- Attack Code A slows down B as the memory utilization increases
- A and B run on the same VM
- Operating System leak is detected. The attack spills over.



The "Experimental Group" for Memory-Attack Experiments:

- Attack Code A runs continuously on VM1 filling up memory
- Computation time is measured for Process B on VM2
- A does not slow down B
- A and B run on different VMs, but same hypervisor, e.g. same cloud
- No hypervisor leak is detected. The attack does not spill over.

(fig. 15)

Analysis and Discussion:

CPU Attacks

As observed in the test results above, the CPU attacks were indeed successful when both the test process and the malware is both launched in the same VM. As seen from figure 7, the time the benign user application took to complete remained at 12 seconds, while the time it took for the malware to complete at the highest priority level lowered to about 6 seconds. Because process A and process B were essentially the same code, they both would normally finish at ~6 seconds each when individually tested at 0 niceness(default niceness). This means that at the highest priority index, the malware code consumed 100% of the CPU's resources,

leaving no processing power to be used by the user application. This is effectively a denial of service attack, which can potentially be used to disrupt a different application's runtime.

Malicious hackers and cyber thieves can potentially use this vulnerability to stifle an application running on a server in a datacenter, which will render the application (and the server) useless.

This can affect many large companies that host large virtualized datacenters. A denial of service attack like this can potentially disable important programs that need to constantly be running in order for the target company to provide its services, for example.

In contrast, no vulnerabilities were found when running the CPU intensive malware on one VM while running a benign application on another. There was no change in the completion time of the user application, so the hypervisor did not allow a leak of computing power to the stressed VM. This reassuring discovery proves that Oracle's Hypervisor is secure when it comes to CPU stress-tests.

RAM Attacks

RAM attacks on one virtual machine were proven effective, since the memory utilization of the virtual machine rose from 59.44% to 85.22%(as seen in figure 15). This increase in memory utilization indicates that it is indeed possible to conduct yet another denial of service attack on one virtual machine with a memory-intensive malware. Attackers can abuse this vulnerability and effectively delay or terminate applications that are running on certain virtual machines. This can be especially damaging to VMs designed to run powerful, resource consuming simulations. An attacker can launch a memory attack like this one and stifle the processes of targeted machines by denying them access to the amount of RAM needed by those machines.

RAM attacks between two VMs, however, was proven impossible. As shown in figure 15, there was no significant change in RAM utilization in the target machine, even when the compromised VM was running the malware at the highest priority level. This shows that Oracle's hypervisor correctly manages the resources of its VMs yet again.

Conclusions:

In this research project, I investigated the robustness of the Oracle VM Virtualbox hypervisor along with individual virtual machines, and my results certainly were notable and surprising. I concluded with great certainty that even if a group of one tenant's VMs are compromised during an attempted security breach, the VMs of other tenants remain unaffected if the attack method was stressing the CPU and RAM resources of the VMs. It was surprisingly easy to compromise a single VM with my CPU and RAM attacks, but the hypervisor prevented a leak of resources between the VMs. These results closely resembled those of previous publications investigating similar attempted breaches into VMs.

Even though I did not find a vulnerability in the actual hypervisor, I certainly found a way to compromise an individual virtual machine. This find is itself a very important vulnerability that is Ubuntu-based.

Future Work:

To expand upon my research, I can investigate different types of storage attacks on virtual machines. Normally, virtualized hard drives are managed by a SAN (storage area network), which manages each VM's accessibility to that hard drive. A study on trying to find vulnerabilities regarding the SAN can be very useful for organizations with large virtualized

datacenters, since it regards the storage infrastructure of all of the virtual machines in said datacenters.

Bibliography

[1] NIST. (2011, January). Guide to Security for Full Virtualization Technologies. Retrieved October 7, 2016, from <http://nvlpubs.nist.gov/nistpubs/Legacy/SP/nistspecialpublication800-125.pdf>

[2] IRACST - International Journal of Computer Science and Information Technology & Security (IJCSITS) Vol. 1, No. 2, December 2011, "Cloud Computing: Security Issues and Research Challenges", from <http://ijcsits.org/papers/Vol1no22011/13vol1no2.pdf>

[3] Oracle Critical Patch Update - January 2016. (2016, January). Retrieved October 07, 2016, from <http://www.oracle.com/technetwork/topics/security/cpujan2016-2367955.html#AppendixOVIR>

[4] Nessus Plugin ID | Tenable Network Security. (2016, January). Retrieved October 07, 2016, from <https://www.tenable.com/plugins/index.php?view=single>

[5] US Census Bureau Public Information Office. (2011, June 29). U.S. Census Bureau Announces Field Management Reforms to Reduce Costs and Enhance Data Quality - Miscellaneous - Newsroom - U.S. Census Bureau. Retrieved October 07, 2016, from <https://www.census.gov/newsroom/releases/archives/miscellaneous/realignment.html>

[6] InfoWorld. (2013, December 18). 4 ways network virtualization improves security. Retrieved October 07, 2016, from <http://www.infoworld.com/article/2609571/networking/4-ways-network-virtualization-improves-security.html>

[7] Exploitdb. (2014, October 8). VirtualBox - 3D Acceleration Virtual Machine Escape (Metasploit). Retrieved October 7, 2016, from <https://www.exploit-db.com/exploits/34334/>

[8] Falcon, F. (2014, June 27). Breaking Out of VirtualBox through 3D Acceleration. Retrieved October 07, 2016, from <https://www.coresecurity.com/corelabs-research/publications/breaking-out-virtualbox-through-3d-acceleration>

[9] Aarseth, R. (2015, September). Security in cloud computing and virtual environments. Retrieved October 7, 2016, from http://bora.uib.no/bitstream/handle/1956/10695/138625252.pdf;jsessionid=F32AF4096B3F0FFEA838BFF5F278879F.bora-uib_worker?sequence=1

[10] Roguine, S. (2014, January 6). 7 Predictions for Virtualization in 2014. Retrieved October 07, 2016, from <http://www.acronis.com/en-us/blog/posts/7-predictions-virtualization-2014>